

Sebastian Alonso Bonilla de la Cruz

sebastian.bonillad@gmail.com | +51 992 625 400 | Lima, Peru | [linkedin.com/in/sebastianbonilladelacruz](https://www.linkedin.com/in/sebastianbonilladelacruz)

PROFILE

Senior Data Engineer with 5+ years designing, developing, and optimizing **ETL/ELT pipelines, Data Lakes, Data Warehouses, and Lakehouse architectures** across cloud environments. Strong proficiency in **Python, PySpark, SQL**, and **Apache Spark** for large-scale data processing, combined with hands-on experience extracting and transforming data from diverse sources and formats (JSON, XML, REST APIs, flat files). Proven track record implementing **data quality methodologies, observability frameworks, and governance standards** throughout ingestion and transformation processes. Deep experience with **Apache Airflow** for pipeline orchestration, **dimensional modeling (star/snowflake schema)**, and **PostgreSQL, Aurora, DynamoDB** across relational and NoSQL environments. Experienced developing and consuming **REST APIs** for data integration, with CI/CD via **Git and GitHub Actions**. Certified **AWS Data Engineer Associate (2026), Microsoft Fabric DP-700, and Azure DP-900**.

WORK EXPERIENCE

SR. DATA ENGINEER | Cheers Health

Jul 2025 – May 2026 Texas, USA (Remote)

- Designed and developed **ETL/ELT pipelines** extracting data from Amazon, Walmart, and TikTok **REST APIs** in JSON format into AWS (S3, Glue, Lambda, Redshift), processing 4M+ daily records with **data quality validation and integrity checks** throughout ingestion and transformation
- Implemented **data quality methodologies** and observability frameworks across all production pipelines; optimized queries and Spark jobs reducing nightly processing windows from 6+ hours to under 45 minutes and cutting operational costs
- Developed **SCD Type 2 CDC patterns** for customer lifecycle analysis; led migration of legacy Xplenty pipelines to modern AWS architecture; enforced CI/CD via **GitHub Actions**
- Built **RAG pipelines** (LangChain, FAISS, OpenAI GPT) integrating LLM capabilities into data workflows; contributed to data science initiatives with clean, well-structured datasets

SR. AWS DATA DEVELOPER | The Hawkers Club

Jan 2024 – May 2025 Denver, USA (Remote)

- Built **Snowflake Data Warehouse** with **dbt and Airflow** (100+ DAGs), designing **star schema dimensional models** and reusable transformation layers processing 800K+ daily transactions from CRM and operational sources
- Developed **Python transformation workflows and REST API integrations** for automated data ingestion; implemented data quality controls reducing pipeline incidents from ~3 daily to near zero
- Designed **data models supporting analytics and BI** reporting on membership economics, churn, and LTV; delivered observability instrumentation and lineage documentation across the platform
- Built **LLM-powered document extraction workflows** (LangChain, RAG, OpenAI GPT) reducing manual validation from ~8 hours to under 1 hour; managed version control and CI/CD via **GitHub Actions**

SR. IMPLEMENTATION CONSULTANT | DigitalCALA

Apr 2021 – Jan 2024 Florida, USA (Remote)

- Architected **Lakehouse solution (Bronze/Silver/Gold)** on Azure Databricks processing 10M+ subscriber records using **PySpark, Apache Spark, dbt, and Azure Data Factory**; integrated multiple enterprise source systems via REST APIs and event streams (Azure Event Hubs)
- Optimized **Spark jobs and SQL queries** from 3 hours to under 40 minutes; reduced data latency from 24 hours to under 2 hours; applied **SCD Type 2, idempotent pipelines, and data governance standards**
- Delivered CI/CD via Azure DevOps and **Git**; mentored junior engineers on scalable design patterns, data quality methodologies, and collaborative development workflows

DATA ANALYST | Pallon

Oct 2020 – Apr 2021 Zurich, Switzerland (Remote)

- Built **Python and SQL ETL pipelines** extracting and transforming data from multiple formats across 50K+ infrastructure inspection datasets, reducing preprocessing time from 6+ hours to under 1 hour
 - Developed anomaly detection pipelines supporting **data science and computer vision** model workflows; contributed to background in ML-adjacent data preparation projects
-

ACADEMIC BACKGROUND

- AWS Certified Data Engineer – Associate, Amazon Web Services, 2026
 - Fabric Data Engineer Associate (DP-700), Microsoft, 2025
 - Azure Data Fundamentals (DP-900), Microsoft, 2025
 - B.Sc. Bioengineering, Universidad de Ingenieria y Tecnologia (UTEC), Lima, Peru, 2018–2023 — Thesis: smart irrigation system with real-time data processing using Azure IoT Hub, Azure Stream Analytics, and Power BI
-

LANGUAGES

- Spanish – Native
- English – Professional working proficiency (C1)